

CSE 314

Report of Assignment 3 on Nachos

Student ID: 1605005

1605006

1605012

Part 1: Threading

Task1:

Implementation:

This task is about implementing a join call in class `nachos.threads.KThread` so that one thread can wait until the completion of another thread. We have made some changes in `join()` and `finish()` methods of the `KThread` class. We have firstly initialized a new variable named `joinThread` to null which will store the thread which called join first on this thread. In the **join() method**, we have checked whether the status of this thread is `statusFinished` or not. If this thread was finished, we returned from the function immediately. Otherwise, we have checked if the `joinThread` variable is null because the result of calling `join()` a second time on the same thread is undefined. If there was no `join()` called on this thread before, then we store the `currentThread` into the `joinThread` variable and block the `currentThread` by `sleep()`. Another thing to be noted is that before `sleep()`, we have disabled the interrupt and after `sleep()`, we have restored the interrupt. At last, in **finish() method**, we have checked if the `currentThread.joinThread` is null or not, if not null, then we make the `joinThread` ready again so that it can finally wake up.

Testing:

We tested our implementation in the `JoinTest` class, where we fork a `PingTest` runnable from the main thread and wait for it to finish by calling `join()` on it. After the `PingTest` has finished, we print a line from the main thread. The fact that this line only gets printed after all the prints in `PingTest` confirms that `join()` is working.

Task2:

Implementation:

This task is about implementing condition variables directly without using any semaphores in class nachos.threads.Condition2. For this task, we have used a LinkedList of KThread named waitQueue. In addition to this, we have used a Lock variable named conditionLock for synchronization. We have made some changes in sleep(), wake() and wakeAll() methods of Condition2 class. To provide atomicity, interrupt has been disabled before doing any change in function and enabled at the moment of departure. In **sleep() method**, conditionLock was released which was acquired by currentThread before. Then we have added this currentThread to the list of waitQueue and currentThread gets blocked by KThread.sleep(). Finally, whenever this thread is woken up from sleep state, conditionLock has been reacquired by that thread again and returns. In **wake() method**, we have checked whether the waitQueue is empty or not and if this queue is not empty, then the first thread of this queue gets woken up by the ready() call which slept on this condition variable before. At last, in **wakeAll() method**, until the waitQueue is not empty, we have called wake() to wake all the sleeping threads.

Testing:

Condition2 was checked in conjunction with the Communicator class. The testing of the communicator class has been explained in the 'Task4' section.

Task3:

Implementation:

This task is about completing the implementation of `nachos.threads.Alarm` class. For this task, we have made changes in `waitUntil(long x)` and `timerInterrupt()` methods in the `Alarm` class. Here, we have used a `priorityQueue` of a custom class named `Entry`. `Entry` stores the waiting threads and their wake time. We have implemented `compareTo(Entry other)` method in `Entry` class so that in `priorityQueue`, all the elements are sorted in ascending order of `wakeTime`. In the **`waitUntil(long x) method`**, we have calculated the `wakeTime` by adding `Machine.Timer.getTime()` to `x`. Then a new object of `Entry` class has been created with this `currentThread` and `wakeTime` and also we have pushed this object to the `priorityQueue`. Finally, `currentThread` gets blocked by `KThread.sleep()` and it will be woken up by `timerInterrupt()` method after passing its `wakeTime`. In **`timerInterrupt() method`**, we have checked if the current time has passed the `wakeTime` of the first entry of the `waitingQueue`, if so we removed it from the queue, and called `ready()` on it to wake it up again. We keep doing this check till we found a entry whose `wakeTime` has still not been passed or the queue has become empty.

Testing:

We have tested this in the `AlarmTest` class. Here, we have forked 3 different `Kthreads`, and called `waitUntil` on all 3 of those for a fixed time. We printed the time before calling `waitUntil`, and also printed the time after they woke up. By checking the difference in sleep and wake time for all of them, we are able to confirm that the `Alarm` class is now working.

Task4:

Implementation:

This task is about implementing class `nachos.threads.Communicator` with operations, `void speak(int word)` and `int listen()` so that a speaker sends and a listener receives msg in a synchronized way. For this task, we have used five `Condition2` variables, one lock and two boolean variables. Firstly, **in `speak(int word)` method**, lock is acquired for mutual exclusion. In `speak()`, a speaker first checks if `isSpeaking` is true, if so, it means that already a speaker is active, so it sleeps on the transferred condition variable. If not, it sets `isSpeaking` to true, to prevent further speaker from interrupting. It then calls `wake` on `speakerReady` variable which would wake a waiting listener if available. It then sleeps on `listenerReady`, waiting for a listener to be active. When a listener wakes the speaker, it transfers the word, and calls `wake()` on the `spoken()` variable, to let the corresponding listener know that it has finished writing. It then waits on the `listened` variable. Finally, when it gets waken by the listener, it sets `isSpeaking` and `isListening` to false, and calls `wake` on `transferred`, to allow a new speaker to be active.

On the other hand, **in `listen()` method**, a listener checks if `isListening` is false, if true it means another listener is active, so it sleeps on the `speakerReady` variable. It also checks if `isSpeaking` is true, if not, it means no speaker is currently active, in that case too, it waits on the `speakerReady` variable. If it passes these two checks, it sets `isListening` to true and calls `wake` on `listenerReady` variable, to wake the currently waiting speaker. It then sleeps on the `spoke` variable and waits for the speaker to speak. After it gets woken up, it stores the `spoken int` and calls `wake` on the `listened` variable to wake the corresponding speaker up. It then returns the received word.

Testing:

We tested both the `Communicator` and `Condition2` in the `CommunicatorTest` class. We forked 3 different speakers and 2 different listeners here. Each speaker speaks from 0 to 5. The fact that a speaker begins speaking only after a listener is available and only one speaker can speak at a time and only one listener can hear

what a speaker just said indicates that the Communicator and Condition2 are working correctly.

Part 2: Multiprogramming

Task1:

Implementation:

This task is about implementing read and write system calls in `nachos.userprog.UserProcess` class so that the user can read the inputs from console and write something in console. For this task, we have added two functions named `handleRead(int fileDescriptor, int vaddr, int size)` and `handleWrite(int fileDescriptor, int vaddr, int size)`. Here, we have used a Lock variable named `ioLock` for atomicity of read and write operation. At the start, when `UserProcess` is created, it initializes its file descriptors by using `openForReading()` and `openForWriting()` of the console of the kernel.

In **`handleRead(int fileDescriptor, int vaddr, int size) method`**, `ioLock` is acquired and then we have checked if `fileDescriptor` is not `stdin` or `vaddr` is less than zero or `size` is less than zero and if any condition is true, we return -1 as error. Otherwise, we have declared a buffer array so that number of bytes copied from console can be stored in that array. Then `readsize` is returned from `openFile.read(buffer,0,size)` method and if `readsize` is greater than zero we have written the bytes of buffer array into virtual memory by calling `writeVirtualMemory(vaddr,buffer,0,readsize)`. Now we have checked if `writeSize` is equal to `readSize` and if they are equal then we return `writeSize`, otherwise we return -1 as error. Finally, `ioLock` is released by the process.

In **`handleWrite(int fileDescriptor, int vaddr, int size) method`**, `ioLock` is acquired and then we have checked if `fileDescriptor` is not `stdout` or `vaddr` is less than zero or `size` is less than zero and if any condition is true, we return -1 as error. Otherwise, we have declared a buffer array so that number of bytes copied from virtual memory can be stored in that buffer. Then `readSize` is returned from `readVirtualMemory(vaddr,buffer,0,readsize)` and if `readsize` is not equal to `size`, we return -1 as error. Otherwise, `writeSize` is returned from `openFile.write(buffer,0,size)` method which defines the number of bytes written on console. At last, if `readSize` and `writeSize` is not equal we return -1 as error and

otherwise we return writeSize. Finally, IoLock is released by the process at the departure.

Testing:

In our shell code, all sorts of printf are working, which indicates that writing to console is being done properly. So, write system call is working. Also, in the forked echo.coff process, we take an integer, and later a string as input, both of which are getting delivered properly. So, the read system call is also working.

Task2:

Implementation:

This task is about implementing support for multiprogramming so that multiple user processes can run at a time. For this task, we have made changes in nachos.userprog.UserProcess class and nachos.userprog.UserKernel class.

In **UserKernel class**, we have used a global LinkedList of free physical pages and a pageLock for atomicity of operation of the list. In **initialize(string[] args) method**, we have initialized the LinkedList by adding all the free pages which is equal to total physical pages(64). We have also shuffled the list for testing so that all pages of a process always cannot be allocated in contiguous block. Then, we have added two functions in UserKernel class named getPage() and freePage(ppn). In **getPage()**, we return a free physical page number from the LinkedList and remove it from this list. If there is no free page in the list, then we return -1 that means insufficient memory. In **freePage(ppn)**, we have added the page back to the LinkedList again which means that this page is free for allocation of new process.

In **UserProcess class**, we have made changes in four functions. We have initialized processID as the number of processes and a pageTable array of TranslationEntry in the constructor of this class. A lock variable is used for the update of processID. In **loadSections() method**, we allocate the memory for this process and load the COFF sections into the memory. Every section has multiple pages. So we have used two for loops, one is iterated through number of sections and another loop is iterated through section.getLength().

Virtual page number is retrieved from each section and a new physical page number is returned from `UserKernel.getNewPage()`. If there is no free physical page, we return false as insufficient free memory and free the already allocated pages. Otherwise, we have added a new `TranslationEntry` to the `pageTable` with `vpn`, free `ppn`, `validBit` as true, `dirtyBit` as false, `usedBit` as false and `readOnlyBit` as `section.isReadOnly()`. Then, we have loaded the physical page into memory by calling `section.loadPage(sectionPageNumber,physicalPageNumber)`. In addition to this, `stackPages` and `argumentPage` are also added to the `PageTable` as `TranslationEntry`. If all the pages of the process are successfully added in the `pageTable`, we return true which means this process will definitely be run. In **`unloadSections() method`**, we have just freed the physical pages which were used for this process by calling `UserKernel.freePage()` and also closed `stdin` and `stdout` fileDescriptors for this process.

In **`readVirtualMemory(int vaddr,byte[] data,int offset,int length) method`**, we have checked if `vaddr` is less than zero or greater than length of main memory and in that case we return 0. Then, `maxAmount` is the number of maximum bytes which can be read from the virtual memory. We have used a for loop which iterates through `maxAmount` and in each iteration, we translate the virtual address to physical address by using the page table. Finally, bytes of memory on this physical address have been stored in the data buffer array by incrementing the amount of bytes transferred and make `usedBit` of that entry as true. At the end of the method, we return the total number of bytes successfully transferred.

In **`writeVirtualMemory(int vaddr,byte[] data,int offset,int length) method`**, like `readVirtualMemory()` function we have added same functionalities in this method. The difference is just when we get physical address by translation, then bytes of data buffer array have been transferred to the memory indexed on that physical address by making `usedBit` and `dirtyBit` of that entry as true. At the end, we return the total number of of bytes successfully transferred to the memory. Before writing to an address, we also first check if that address belongs to a read only page, if so, we return without writing any further.

Testing:

As the read and write system calls use the readVirtualMemory and writeVirtualMemory functions, and the fact that all the reads and writes in our shell program is still working, indicates that the virtual memory allocation system is working. Also, as we are able to create several processes in turns, it means that the exiting processes are also freeing their memory properly.

Task3:

Implementation: This task is about implementing three system calls which are exec, join and exit in nachos.userprog.UserProcess class. For this task, we have added three functions in the class named handleExec(int nameVAddress,int argc,int argVAddress), handleJoin(int pid,int statusVAddress) and handleExit(int status).

We have added some new data structures such as a List of UserProcess named childProcesses, parentProcessID (int), activeProcessCount (int), exitStatus (int) and didExitNormally (boolean) for implementation of these system calls.

In **handleExec(int nameVAddress,int argc,int argVAddress) method**, if VAddress, argc or argvVAddress is less than zero, we return -1. Otherwise, we have read file name from virtual memory by calling readVirtualMemorystring(nameVAddress,maxFileSize). If this file name does not end with “.coff”, we return -1. Otherwise, we have made a argV array of String which contains all the arguments passed to the new process. After that, new userProcess has been initialized and we have checked if new process can execute the coff file. And if there is no error in the execution of that coff file then child process set its parentprocessID as this processID and also we have added the new child process to the childProcessesList of current process. Finally, we return the processID of child process.

In **handleExit(int status) method**, all the children of this process disown its parent by setting parentProcessID as -1. Then, all the sections will be unloaded and descriptors will be closed for this process by calling unloadSections(). Now, we decrement the activeProcessCount as this process is no longer to be run. If this count is zero that means this is the last exiting process, then we call Kernel.kernel.terminate() directly. Otherwise, the status has been stored in exitStatus variable which is passed to parent process and user thread of this current process has been destroyed by calling UThread.finish().

In **handleJoin(int pid, int statusVAddress) method**, we have iterated through each child process and in each iteration we have checked if pid is equal to childProcessID and if that is true we call child.processThread.join() for joining on child process. Hence, currentProcess gets blocked then and waits until completion of that child process. Whenever parent process is ready to run, we have removed this child process from the List so that parent process cannot join on the same child process again in future. Now, exitStatus of child process has been passed to parent process and this exitStatus bytes have been written into virtual memory starting at statusVAddress. At last, we have checked the value of didExitNormally. If it is false, it means that exiting of the process occurred due to an unhandled exception and so we return 0. So, in handleException(int cause) method, we call handleExit(-1) in the case of unexpected Exception. Otherwise, if the process exited normally, we return 1 as a successful join. If pid is not equal to any processId of any child, we return -1 as invalid join.

In **handleHalt() method**, we have added extra check whether the current processID is zero or not. Because only the root process can invoke the halt system call. So, if processID is not 0, we return -1. Otherwise, Machine.halt() is called.

Testing:

In our shell code (myPgr.c), we have used exec to create several child processes. All of those processes are working. We have also called join on these processes, and the parent process waits for the child processes to end. And if we call join on processes with invalid ids, join returns -1. Which means that join is working. Also, the exitStatus of the children is being written properly, which has been verified by printing them after they join. We have also checked for exit of a process due to error by calling exec on “bad.coff”, in which case, the OS does not crash, and the parent process gets -1 as the exitStatus of that process. We have also tried halting from non-root processes, which fail to halt the machine, but the root can halt the machine successfully. So halt is also working correctly.